

IIT Madras BSc Degree

Copyright and terms of use

IIT Madras is the sole owner of the content available in this portal - onlinedegree.iitm.ac.in and the content is copyrighted to IIT Madras.

- Learners may download copyrighted material for their use for the purpose of the online program only.
- Except as otherwise expressly permitted under copyright law, no use other than for the purpose of the online program is permitted.
- No copying, redistribution, retransmission, publication or exploitation, commercial or otherwise of material will be permitted without the express permission of IIT Madras.
- Learner acknowledges that he/she does not acquire any ownership rights by downloading copyrighted material.
- Learners may not modify, publish, transmit, participate in the transfer or sale, create derivative works, or in any way exploit, any of the content, in whole or in part.

Models

Persistent Storage

- Example: Grades
- Need for persistent storage
- Requirements

Example: Gradebook

- Students: ID, name, address, ...
- Courses: ID, name, department, year, ...
- StudentCourse Relationship: which students are registered for which courses

Gradebook

	A	B
1	Name	IDNumber
2	Sunil Shashi	MAD001
3	Chetana Anantha	MAD002
4	Madhur Prakash	MAD003
5	Nihal Surya	MAD004
6	Shweta Lalita	MAD005
7	Raghu Balwinder	MAD006
8	Gulshan Kuldeep	MAD007
9	Kishan Shrivatsa	MAD008
10	Purnima Sunil	MAD009
11	Nikitha Madhavi	MAD010
12	Lilavati Prabhakar	MAD011
13	Rama Yamuna	MAD012

	A	B
1	CourseID	Name
2	EE1001	Introduction to Electrical Engineering
3	AM1100	Engineering Mechanics
4	MA1020	Functions of Several Variables
5	ME1100	Thermodynamics
6	BT1010	Life Sciences

Spreadsheets

- Arbitrary data organized into Rows and Columns
- Operations defined on Cells or Ranges
- Multiple inter-linked sheets within single spreadsheet

Any kind of **tabular** data - expressed in tables

Relationships

- Student - Course?

Relationships

- Student - Course?
- Separate entry with full details - student name, ID, address, course ID, name, department etc?

Relationships

- Student - Course?
- Separate entry with full details - student name, ID, address, course ID, name, department etc?
 - Redundant

Relationships

- Student - Course?
- Separate entry with full details - student name, ID, address, course ID, name, department etc?
 - Redundant
- Separate table “joining” students with courses
 - Only ID specified!
 - Relation specified with “Keys”

	A	B	C
4	MAD001	BT1010	78
5	MAD002	EE1001	30
6	MAD005	EE1001	68
7	MAD009	AM1100	62
8	MAD012	AM1100	77
9	MAD007	BT1010	41
10	MAD001	MA1020	56

Questions

- How should the underlying data be stored?
 - Can it be made persistent - survive server restart?
- How should the relations be represented?
- Structured ways to represent, manipulate data?

Storage

Mechanisms for persistent storage

In memory data structures

```
names = ['Alice', 'Bob', 'Charlie']
courses = ['Introduction to EE', 'Applied Mech', 'Calculus']
rels = [('Alice', 'Introduction to EE'),
        ('Bob', 'Calculus'),
        ('Alice', 'Calculus'),
        ('Charlie', 'Applied Mech')]
```

In memory data structures

```
names = ['Alice', 'Bob', 'Charlie']
courses = ['Introduction to EE', 'Applied Mech', 'Calculus']
rels = [('Alice', 'Introduction to EE'),
        ('Bob', 'Calculus'),
        ('Alice', 'Calculus'),
        ('Charlie', 'Applied Mech')]
```

- Error prone - easy to make mistakes in entry or referencing
- Does not scale
- Duplicate names?

In memory data structures - Keys

```
names = {0: 'Alice', 1: 'Bob', 2: 'Charlie'}  
courses = {0: 'Introduction to EE',  
            1: 'Applied Mech',  
            2: 'Calculus'}  
rels = [(0, 0),  
        (1, 2),  
        (0, 2),  
        (2, 1)]
```

In memory data structures - Keys

```
names = {0: 'Alice', 1: 'Bob', 2: 'Charlie'}  
courses = {0: 'Introduction to EE',  
            1: 'Applied Mech',  
            2: 'Calculus'}  
rels = [(0, 0),  
        (1, 2),  
        (0, 2),  
        (2, 1)]
```

- Data entry errors less likely
- Duplicates not a problem - Unique **Key**

Objects

```
class Student:
    idnext = 0 # Class variable
    def __init__(self, name):
        self.name = name
        self.id = idnext
        idnext = idnext + 1
```

- Auto-initialize ID to ensure unique
- Functions to set/get values

Objects

```
class Student:
    idnext = 0 # Class variable
    def __init__(self, name, hostel):
        self.name = name
        self.id = idnext
        self.hostel = hostel
        idnext = idnext + 1
```

- Add a new field to object easily

Persistence?

- In memory data structures lost when server shut down or restarted
- Save to disk? Structured data?
 - Python Pickle and similar modules
 - CSV - comma separated values
 - TSV - tab separated values
- Essentially same as spreadsheets: limited flexibility

Spreadsheet

- Naturally represent tabular data
- Extension, adding fields easy
- Separate sheet for relationships

Spreadsheet

- Naturally represent tabular data
- Extension, adding fields easy
- Separate sheet for relationships

Problems:

- Lookups, cross-referencing harder than dedicated database
- Stored procedures - limited functionality
- Atomic operations - no clear definition

Relational Databases - SQL

- From IBM ~ 1970s
- Data stored in Tabular format:
 - Columns of tables: fields (name, address, department, ...)
 - Rows of tables: individual entries (student1, student2, ...)

Unstructured databases - NoSQL

- Easily add/change fields
- Arbitrary data
- NoSQL
 - MongoDB
 - CouchDB
 -
- Flexible, but potential loss of validation

Relations

Spreadsheet

	A	B	C
4	MAD001	BT1010	78
5	MAD002	EE1001	30
6	MAD005	EE1001	68
7	MAD009	AM1100	62
8	MAD012	AM1100	77
9	MAD007	BT1010	41
10	MAD001	MA1020	56

	A	B
1	CourseID	Name
2	EE1001	Introduction to Electrical Engineering
3	AM1100	Engineering Mechanics
4	MA1020	Functions of Several Variables
5	ME1100	Thermodynamics
6	BT1010	Life Sciences



	A	B
1	Name	IDNumber
2	Sunil Shashi	MAD001
3	Chetana Anantha	MAD002
4	Madhur Prakash	MAD003
5	Nihal Surya	MAD004
6	Shweta Lalita	MAD005
7	Raghu Balwinder	MAD006
8	Gulshan Kuldeep	MAD007
9	Kishan Shrivatsa	MAD008
10	Purnima Sunil	MAD009
11	Nikitha Madhavi	MAD010
12	Lilavati Prabhakar	MAD011
13	Rama Yamuna	MAD012

Relationship types

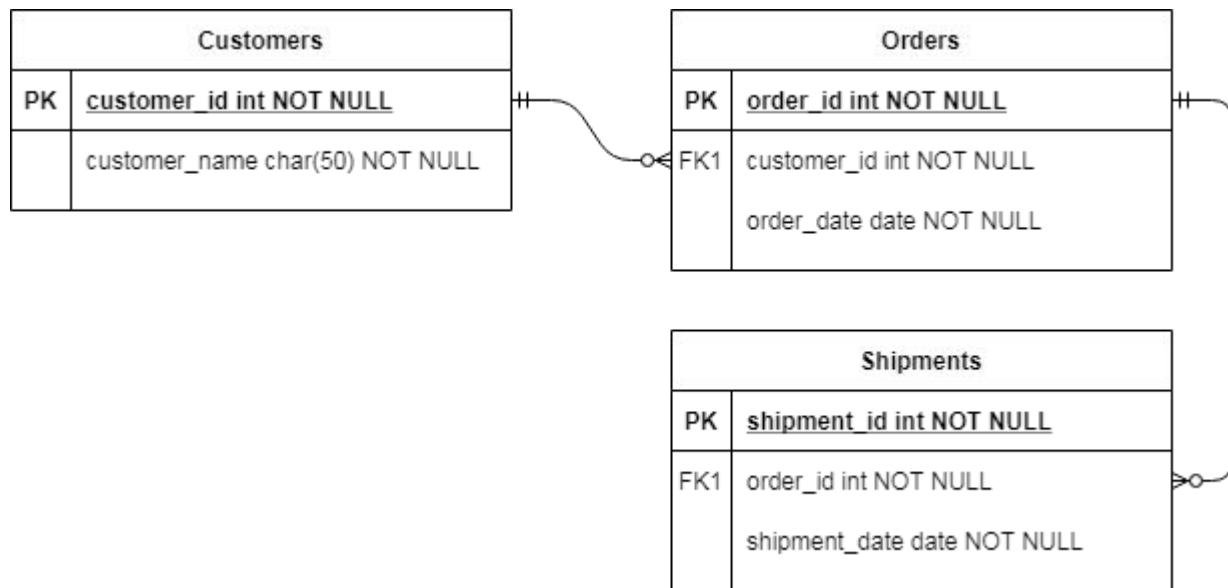
- One-to-one:
 - One student has one roll number
 - One roll number uniquely identifies one student
 - Example: assign unique message-ID to each email in Inbox
- One-to-many (many-to-one):
 - one student stays in only one hostel
 - one hostel has many students
 - Example: save emails in folders - one email is in only one folder
- Many-to-many:
 - one student can register for many courses
 - one course can have many students
 - Example: assign labels to emails - one email can have many labels and vice versa

Diagrams




- Entity-Relationship
- UML
- Class relation
- ...



Entity-Relationship Diagram



ER Examples



Tool: Draw.io - <https://app.diagrams.net/>

SQL

Relational Databases

- From IBM ~ 1970s
- Data stored in Tabular format:
 - Columns of tables: fields (name, address, department, ...)
 - Rows of tables: individual entries (student1, student2, ...)
- Key: unique way of accessing a given row

Relational Databases

- From IBM ~ 1970s
- Data stored in Tabular format:
 - Columns of tables: fields (name, address, department, ...)
 - Rows of tables: individual entries (student1, student2, ...)
- Key: unique way of accessing a given row
 - **Primary key**: important for fast access on large databases

Relational Databases

- From IBM ~ 1970s
- Data stored in Tabular format:
 - Columns of tables: fields (name, address, department, ...)
 - Rows of tables: individual entries (student1, student2, ...)
- Key: unique way of accessing a given row
 - **Primary key**: important for fast access on large databases
 - **Foreign key**: connect to a different table - **Relationships**

Queries

- Retrieve data from the database:

eg. “Find students with name beginning with A”

“Find all courses offered in 2021”

Structured Query Language (SQL)

- English - like, but structured
- Quite verbose
- Specific mathematical operations:
 - Inner Join
 - Outer Join

Example: Inner Join

Name	IDNumber	hostelID
Sunil Shashi	MAD001	1
Chetana Anantha	MAD002	2
Madhur Prakash	MAD003	2
Nihal Surya	MAD004	3
Shweta Lalita	MAD005	2
Raghu Balwinder	MAD006	3
Gulshan Kuldeep	MAD007	1
Kishan Shrivatsa	MAD008	1
Purnima Sunil	MAD009	2
Nikitha Madhavi	MAD010	1
Lilavati Prabhakar	MAD011	3
Rama Yamuna	MAD012	3

ID	Name	Capacity
1	Jamuna	300
2	Ganga	300
3	Brahmaputra	500

Student - Hostel mapping

```
select Students.name, Hostels.name  
  from Students  
 inner join Hostels  
on Students.hostelID = Hostels.ID
```

Student - Hostel mapping

```
select Students.name, Hostels.name  
  from Students  
 inner join Hostels  
  on Students.hostelID = Hostels.ID
```

Sunil Shashi, Jamuna

Chetana Anantha, Ganga

Cartesian Product

- N entries in table 1
- M entries in table 2
- M x N combinations - filter on them

Powerful SQL queries can be constructed

Example: find all students in Calculus

- Find ID number for course
- Look up StudentsCourses table to find all entries with this course ID
- Look up Students to find names of students with these IDs

Example: find all students in Calculus

- Find ID number for course
- Look up StudentsCourses table to find all entries with this course ID
- Look up Students to find names of students with these IDs

```
SELECT s.name  
FROM Students s
```

Example: find all students in Calculus

- Find ID number for course
- Look up StudentsCourses table to find all entries with this course ID
- Look up Students to find names of students with these IDs

```
SELECT s.name  
FROM Students s  
JOIN StudentsCourses sc ON s.IDNumber = sc.studentID
```

Example: find all students in Calculus

- Find ID number for course
- Look up StudentsCourses table to find all entries with this course ID
- Look up Students to find names of students with these IDs

```
SELECT s.name  
FROM Students s  
JOIN StudentsCourses sc ON s.IDNumber = sc.studentID  
JOIN Courses c ON c.ID = sc.courseID
```

Example: find all students in Calculus

- Find ID number for course
- Look up StudentsCourses table to find all entries with this course ID
- Look up Students to find names of students with these IDs

```
SELECT s.name
FROM Students s
JOIN StudentsCourses sc ON s.IDNumber = sc.studentID
JOIN Courses c ON c.ID = sc.courseID
WHERE c.name='Calculus'
```

Summary

- Models - persistent data storage
- Mechanisms:
 - CSV, Spreadsheets, SQL, NoSQL
- Entities and Relationships
 - Different ways of representing

No details on display, views, or what kind of updates permitted